

On the Economic and Governance Mechanisms for the Agentic Web: A Global South Perspective

Diogo Cortiz
NIC.br / PUC-SP
São Paulo, Brazil
diogo@nic.br

Beatriz Rocha
NIC.br
São Paulo, Brazil
biarocho@nic.br

Vagner Figueredo de Santana
NIC.br
São Paulo, Brazil
vsantana@nic.br

Henrique S. Xavier
NIC.br
São Paulo, Brazil
hxavier@nic.br

Jacqueline Custodio
NIC.br
São Paulo, Brazil
custodiosilva@nic.br

Abstract

The Agentic Web signals a transition from a Web organized around documents to one increasingly mediated by autonomous AI agents. This shift expands possibilities for interaction and automation, but it also disrupts the technical and economic foundations of the Web. Drawing on a multi-stakeholder workshop held in Brazil with 28 representatives from major technology companies, media groups, newspapers, academia, and government, this position paper examines these changes from a Global South situated knowledge perspective. We highlight three central tensions: the lack of standardized mechanisms to control AI-bots access to content; the growing mismatch between agent-based content consumption and current digital business models; and governance gaps that might amplify extractive dynamics. Current responses, such as broad bot blocking and legal disputes, show a deeper misalignment among the actors involved in the development of agentic Web. We argue that addressing this misalignment requires rebalancing incentives through transparent standards, inclusive governance, and multi-stakeholder coordination, so that the Agentic Web can evolve as an open and sustainable infrastructure.

CCS Concepts

• Information systems → World Wide Web; • Social and professional topics → Computing / technology policy.

Keywords

Agentic Web, Global South, Artificial Intelligence, Governance, Regulation

ACM Reference Format:

Diogo Cortiz, Beatriz Rocha, Vagner Figueredo de Santana, Henrique S. Xavier, and Jacqueline Custodio. 2026. On the Economic and Governance Mechanisms for the Agentic Web: A Global South Perspective. In *Companion Proceedings of the ACM Web Conference 2026 (WWW Companion '26)*, April 13–17, 2026, Dubai, United Arab Emirates. ACM, New York, NY, USA, 4 pages. <https://doi.org/10.1145/3774905.3795089>



This work is licensed under a Creative Commons Attribution 4.0 International License. *WWW Companion '26, Dubai, United Arab Emirates*
© 2026 Copyright held by the owner/author(s).
ACM ISBN 979-8-4007-2308-7/2026/04
<https://doi.org/10.1145/3774905.3795089>

1 Introduction

With the advent of Generative Artificial Intelligence (AI) and Large Language Models (LLMs), content previously available only on the Web is now embedded in or presented through these models. Moreover, LLM-based agents are now being used to perform more complex tasks on the Web on behalf of users. The emergence of the "Agentic Web" promises a paradigm shift from a Web of documents to a Web of autonomous entities [14].

As LLM-based agents mediate human interaction with information, the foundational structures of the internet – technical, economic, and cultural – are being redefined. While much of the discourse focuses on the optimization of agentic workflows [10, 13, 15, 23], there is a need to evaluate how these transformations would socially and economically impact diverse ecosystems, particularly in the Global South [8], where concepts such as *digital sovereignty*¹ and *digital colonialism*² are key for ongoing debates [19]. Literature presents proposals for restructuring AI governance frameworks to be more inclusive [20], including perspectives of communities most directly impacted economically [16], and ensuring agency to countries in the Global South instead of spectator, victim roles [12].

By recognizing these complexities, we organized a multi-stakeholder in-person workshop called "Agentic Web: the impact of AI on the Web" in São Paulo, SP, Brazil, to discuss how AI is reshaping the Web, its business models, and social practices associated with it. The event brought together 28 stakeholders including prominent academics, members of the Brazilian Internet Steering Committee (CGL.br), and representatives from Big Techs, media groups, leading newspapers, the Brazilian federal government, and Brazilian regulatory agencies. This diverse representation ensured a plural debate grounded on concrete technical, political, and economical realities. This paper builds on top of results from the workshop's report [22] towards three timely and central topics identified in the discussions:

- (1) Technical challenges for controlling AI-originating access;
- (2) Challenges to current business models; and
- (3) Governance and regulation from a Global South perspective.

¹Digital sovereignty can be defined as the ability a country has in terms of controlling its own digital destiny, be it in terms of data, hardware and software that it relies on and creates [11].

²Digital colonialism can be defined as the power dynamics placing Global South countries as captive consumers of technologies, data providers, and cheap labor, with no sovereignty about the future presented as inevitable [19].

This paper is structured as follows: section 2 details the methodology followed in the workshop, while section 3 presents reflections on selected technical, economical, and governance challenges identified in the workshop. Finally, in section 4, we conclude and propose steps for a more balanced Web ecosystem in the context of the emergence of a “new form of dependency arising from digital-era capitalism” [6], analyzed from the perspective of a specific Global South country, namely, Brazil.

2 Workshop Methodology

To promote open discussion and mutual learning, the workshop followed the Chatham House Rule [7], allowing participants to speak as individuals and share candid, sensitive insights while guaranteeing anonymity. Hence, details allowing one to identify participants or organizations cannot be provided. This approach supported a dialogue that extended beyond official corporate positions and contributed to identifying tensions between content ownership and AI consumption.

The 28 participants were divided into four discussion subgroups. These groups were designed to foster pluralistic discussions by bringing together individuals from different organizations, with diverse roles and skill sets. The workshop discussions were organized around three main topics:

- (1) Technologies and standards;
- (2) User behavior and social impacts; and
- (3) Business models and economic impacts.

For each topic, the subgroups were invited to discuss them from three main perspectives, according to the Theory of Change followed in the workshop [2]:

- (1) **Diagnosis:** Participants assessed the current situation, identifying key challenges and opportunities.
- (2) **Ongoing actions:** Participants listed the actions that are currently being undertaken by their organizations.
- (3) **Future strategies:** Participants engaged in a brainstorming session to proposed future actions for the problems.

Following the subgroup discussions, participants were asked to write short summaries of their conclusions on a flip-chart and present them to the other groups. The workshop concluded with a plenary discussion involving all participants, during which areas of consensus and disagreement were highlighted. A facilitator then prepared a report describing and synthesizing the workshop findings, based on Semantic Thematic Analysis. The full report was shared with the participants for comments and it subsequently made available online [22].

3 Discussion

The workshop’s report presents multiple findings. Many of the current challenges identified are related to the deterioration of current content provider business models, such as:

- the lack of standards for specifying content copyright and licenses;
- the cost publishers bare for serving requests from bots that contribute little in return; and
- the difficulty for small, regional publishers to strike beneficial licensing deals with large AI companies.

The participants also mentioned the dangers of increased misinformation caused by AI hallucinations that could more severely impact the Global South due to the scarcity of representative training and grounding data [20].

Most of the current responses to these challenges are defensive, such as blocking bot access to websites and entering in legal battles against AI companies over copyright infringement. This finding indicates that the advancements brought by the Agentic Web are not properly aligned with the needs and interests of other Web actors, and that such an alignment could turn resistance into fruitful collaboration. As steps envisioned for improving the Web ecosystem, the participants mentioned initiatives such as the promotion of media and AI literacy, the creation and public financing of digital commons [9], the establishment of public policies to treat power imbalances, and incentives for multi-stakeholder and globally diverse models of governance.

In this context, this paper deep dives into three issues that were highlighted in the workshop:

- technical challenges for controlling bot access;
- obsolescence of current digital business models; and
- governance and regulation from a Global South perspective.

3.1 Technical Challenges for Controlling Bot Access

A primary finding of the workshop is the lack of standardized mechanisms for bot permission control, a deficit that could impact the foundational architecture of the Web. The Web was engineered as a decentralized network of hyperlinks connecting documents designed to facilitate the free flow of information and to encourage the sharing of content [4]. On top of it, business models were built based on monetization of visibility and referral traffic [21]. However, the rise of the Agentic Web is fundamentally altering this perspective, as the Web is transitioning from a Web of “links” to a Web of “data extraction,” in which users never reach a publisher’s website despite their content being ultimately consumed by users.

Within this new paradigm, publishers and media organizations currently struggle to distinguish between traditional search engine crawlers (which drive traffic) and AI-specific bots (which often ingest content for model pre-training or inference).

This lack of transparency creates an unsustainable burden on the infrastructure. For example, the Wikimedia Foundation reported that multimedia downloads grew 50% from January 2024 to April 2025, mainly due to automated programs that collect images for AI training [18]. Such practice *significantly increases the operational and bandwidth costs of an organization dedicated to keeping content freely accessible to humans*.

This agent-mediated load is perceived as a systemic threat. Participants argued that current tools (like *robots.txt*) are insufficient for a publisher to specify whether content can be:

- (1) Indexed for traditional search results;
- (2) Used as training data for LLMs; or
- (3) Accessed in real-time for AI inference.

The absence of such standards is forcing publishers to implement aggressive, total blocks on all automated traffic. This “defensive walling” of content threatens to fragment and weaken the Web, moving it away from its open roots and toward a series of

disconnected, private silos. While bodies like the IETF (Internet Engineering Task Force) and W3C (World Wide Web Consortium) are currently debating these issues, the workshop concluded that the *pace of standard-setting is currently lagging behind the speed of AI deployment and adoption*.

3.2 Obsolescence of Current Digital Business Models

The economic discussion highlighted a fundamental “unfair” relationship in current digital business models, especially regarding AI inference. Beyond legal debates over training data copyright, participants highlighted an immediate threat: “real-time inference.” Chatbots collect content in real-time and provide direct summaries to users. This effectively *blocks referral traffic to original creators*, who subsequently lose advertising revenue and subscription opportunities. In addition to direct revenue, participants cited the *loss of authorship* and “authority” in the eyes of the user, as the AI becomes the primary interface, obscuring the original source of the information.

There is also a *data obfuscation*, as the shift from human visitors to AI agents has rendered traditional Web analytics obsolete. *Publishers can no longer track user behavior, interests, or demographics when an agent acts as an intermediary*. The web is transitioning from being rendered for human eyes to be rendered to agents. An evidence of this transition is that 65% of Wikipedia’s most resource-intensive traffic now originates from automated bots [18]. As a consequence, this paradigm shift is now disrupting traditional business models depending on analytics and human behavior. Power balance is key for development and governance of more responsible technologies [3]. However, as this scenario unfolds, power imbalances in the Web ecosystem will potentially exacerbate, as a handful of companies are centralizing access to state-of-the-art models, agents, and agentic browsers, and such imbalances could harm the very content production process required by AI agents.

In dialogues between publishers and technology companies, two paths have emerged as potential solutions, as discussed in a recent W3C report [14]:

- **Micro-transactions:** Implementing technical standards for micro-payments to compensate for specific content hits by agents.
- **Analytics APIs:** Developing new APIs where AI agents can “report back” anonymized analytics data to the original platforms, allowing services to understand their reach in an agent-mediated world.

3.3 Governance and Regulation

The transition to an Agentic Web is often framed in technical circles as a neutral progression toward greater efficiency. However, from the perspective of the Global South, this evolution is a geopolitical battlefield.

The infrastructure, capital, and data processing power required to sustain billions of AI agents are currently concentrated within a handful of jurisdictions, primarily the United States and China, which participate in data colonialism despite their divergent political-economic regimes as liberal and state-led market societies, respectively [17]. Based on that, participants highlighted that this

centralization threatens to solidify a new form of “digital extractivism” [19].

A recurring concern in the workshop was the risk of Global South nations being relegated to the role of “data providers”, as raw material suppliers for models that are then sold back as finished services.

The primary challenge for these countries is to move up the value chain. This involves not only hosting agents but ensuring that the “intelligence” mediating the local Web understands local legal contexts, social norms, and economic priorities.

The group discussed two paths for engaging with this issue.

- (1) The first involves the creation of national policies to promote research, development, and the adoption of technologies and standards.
- (2) The second relates to the discussion of regulatory strategies aimed at bringing a better balance of power and responsibilities.

The Brazilian Case: A Regulatory Framework

Brazil can now be positioned among the pioneers in AI regulation due to Senate Bill 2338/2023 [5]. Recently approved by the Senate and currently under deliberation in the Chamber of Deputies of Brazil, this bill serves as a cornerstone for “sovereignty-by-design.”

While inspired by the European Union’s AI Act³, the Brazilian framework adapts its risk-based approach to the specificities of a peripheral economic position. It is expected to prioritize the protection of fundamental rights without stifling the growth of local AI startups.

The bill seeks to establish a regulatory environment in which innovation is coupled with mandatory transparency. For the Agentic Web, this means that agents operating within Brazilian jurisdiction must be accountable, traceable, and explainable.

Sovereignty and Multi-stakeholder Governance

Workshop participants discussed the principles of the Brazilian Internet Steering Committee (CGI.br) as an inspirational model. CGI.br is often cited as an international successful example of the multistakeholder model of governance, bringing together government, the private sector, civil society, the technical community, and academia in a single institutional arrangement [1]. Created in the early days of the Brazilian Internet, CGI.br operates through shared decision-making, consensus-building, and transparency, ensuring that no single actor dominates the agenda. In addition, CGI.br seeks to ensure a balanced representation of Brazil’s five geographic regions in all its initiatives, including meetings, committees, event locations and grants programs.

Based on this experience, participants argued that the Agentic Web’s foundational protocols and ecosystem must be developed and adopted through an open, transparent process involving the private sector, academia, government, and civil society and that includes a fair share of representatives from the Global South. This could *prevent the emergence of private monopolies and ensures that*

³<https://www.europarl.europa.eu/topics/en/article/20230601STO93804/eu-ai-act-first-regulation-on-artificial-intelligence>

the Web remains a public good suited to several needs, even as it becomes increasingly mediated by autonomous agents.

4 Conclusion

The Agentic Web is emerging in a largely uncoordinated manner, driven in particular by the economic power of large corporations. While the implications for different stakeholders are not yet fully understood, early evidences generated by our workshop suggests potential negative impact for some actors, particularly content producers and public-interest platforms.

As a result, these actors are adopting defensive strategies such as blocking automated access and pursuing legal disputes. These frictions show conflicting interests among the different parties involved in the development of the Agentic Web ecosystem and may lead to unintended consequences, such as technical fragmentation of the Web and economic imbalances among different actors.

This process is negative not only for the Agentic Web itself, but for the society as a whole. A Web shaped by access barriers, content silos, and losing diverse content producers weakens the very ecosystem that sustains AI systems. *If trust erodes and collaboration is replaced by extraction, the development of agentic technologies would become fragile and constrained.*

This paper argues for an alternative approach that aligns the Agentic Web with the needs of other Web actors. The proposed solutions focus on recognizing these needs and rebalancing incentives to support a more collaborative and sustainable environment for developing emerging technologies and business models. Rather than defensive resistance, this alignment can help turn current tensions into more equitable arrangements.

In terms of limitations of this position paper, although the workshop that this work is based of involved stakeholders from multiple sectors and organizations, it may not be a representative sample for the Brazil nor Global South.

From a Global South perspective, this harmonization is fundamental so that actors can challenge exclusionary governance mechanisms while proposing alternative ones, provide legitimate expertise in localization of risks [20], and maintain and improve businesses and other relevant activities on the Web. Without it, the Agentic Web risks reinforcing existing power asymmetries and consolidating extractive dynamics. By contrast, inclusive governance, transparent standards, and multi-stakeholder approaches - particularly those that actively incorporate perspectives, interests, and knowledge from the Global South - can help ensure that the Agentic Web evolves as a shared infrastructure for supporting innovation while preserving the diversity and openness of the Web's core principles.

References

- [1] Kimberly Anastácio. 2015. *Brazil's Approach to Multistakeholderism: Multi-participation in the Brazilian Internet Steering Committee (CGL.br)*. Technical Report. Berkman Center for Internet & Society at Harvard University. https://cyber.harvard.edu/sites/cyber.harvard.edu/files/Publish_Kimberly%20Anastacio.pdf A case study examining stakeholders in the Brazilian Internet Steering Committee (CGL.br).
- [2] Andrea A Anderson. 2004. Theory of Change. *As a Tool for Strategic Planning. A Report on Early Experiences. The Aspen Institute-Roundtable on Community Change* (2004).
- [3] Juana Catalina Becerra Sandoval, Vagner Figueredo de santana, Sara Berger, Lauren Thomas Quigley, and Stacy Hobson. 2025. A Responsible and Inclusive Technology Framework for Attending to Business-to-Business Contexts. *ACM J. Responsib. Comput.* 2, 1, Article 3 (April 2025), 24 pages. doi:10.1145/3712310
- [4] Tim Berners-Lee and Mark Fischetti. 2000. *Weaving the Web: The Original Design and Ultimate Destiny of the World Wide Web*. Harper Business, New York, NY. First edition.
- [5] Brazilian Senate and Senator Rodrigo Pacheco. 2023. Projeto de Lei nº 2338, de 2023: Dispõe sobre o uso da Inteligência Artificial. Atividade Legislativa — Senado Federal. <https://www25.senado.leg.br/web/atividade/materias/-/materia/157233>
- [6] João Francisco Cassino. 2021. O Sul global e os desafios pós-coloniais na era digital. In *Colonialismo de dados: como opera a trincheira algorítmica na guerra neoliberal*. Autonomia Literária.
- [7] Chatham House. 2025. Chatham House Rule. <https://www.chathamhouse.org/about-us/chatham-house-rule>.
- [8] Nour Dados and Raewyn Connell. 2012. The Global South. *Contexts* 11, 1 (2012), 12–13. doi:10.1177/1536504212436479
- [9] Mélanie Dulong de Rosnay and Felix Stalder. 2020. Digital commons. *Internet Policy Review* 9, 4 (2020). doi:10.14763/2020.4.1530
- [10] Abul Ehtesham, Aditi Singh, Gaurav Kumar Gupta, and Saket Kumar. 2025. A survey of agent interoperability protocols: Model Context Protocol (MCP), Agent Communication Protocol (ACP), Agent-to-Agent Protocol (A2A), and Agent Network Protocol (ANP). arXiv:2505.02279 [cs.AI] <https://arxiv.org/abs/2505.02279>
- [11] Sean Fleming. 2025. *What is digital sovereignty and how are countries approaching it?* World Economic Forum. <https://www.weforum.org/stories/2025/01/europe-digital-sovereignty/>
- [12] Eugenio Garcia. 2021. The International Governance of AI: Where Is the Global South? *The Good AI* (2021).
- [13] Idan Habler, Ken Huang, Vineeth Sai Narajala, and Prashant Kulkarni. 2025. Building A Secure Agentic AI Application Leveraging A2A Protocol. arXiv:2504.16902 [cs.CR] <https://arxiv.org/abs/2504.16902>
- [14] Dominique Hazaël-Massieux. 2024. *AI & the Web: Understanding and Managing the Impact of Machine Learning Models on the Web*. Report W3C Report. World Wide Web Consortium (W3C). <https://www.w3.org/reports/ai-web-impact/> Draft originally maintained in GitHub repository w3c/ai-web-impact.
- [15] Xinyi Hou, Yanjie Zhao, Shenao Wang, and Haoyu Wang. 2025. Model Context Protocol (MCP): Landscape, Security Threats, and Future Research Directions. arXiv:2503.23278 [cs.CR] <https://arxiv.org/abs/2503.23278>
- [16] Amba Kak. 2020. "The Global South is everywhere, but also always somewhere": National Policy Narratives and AI Justice. In *Proceedings of the AAAI/ACM Conference on AI, Ethics, and Society* (New York, NY, USA) (AIES '20). Association for Computing Machinery, New York, NY, USA, 307–312. doi:10.1145/3375627.3375859
- [17] Sebastián Lehuédé. 2024. An alternative planetary future? Digital sovereignty frameworks and the decolonial option. *Big Data and Society* (2024), 1–13.
- [18] Birgit Mueller, Chris Danis, Giuseppe Lavagetto, and Wikimedia Foundation. 2025. *How Crawlers Impact the Operations of the Wikimedia Projects*. <https://diff.wikimedia.org/2025/04/01/how-crawlers-impact-the-operations-of-the-wikimedia-projects/>
- [19] Mauro Oliveira. 2025. *Soberania Digital: Colonização e Letramento*. Clube de Autores, Joinville, SC, Brazil. <https://clubedeautores.com.br/livro/soberania-digital> 1st edition, 240 pages.
- [20] Marie-Therese Png. 2022. At the tensions of south and north: Critical roles of global south stakeholders in AI governance. In *Proceedings of the 2022 ACM Conference on Fairness, Accountability, and Transparency*. 1434–1445.
- [21] Jeffrey F Rayport. 1999. The truth about Internet business models. *Strategy and Business* (1999), 5–7.
- [22] Beatriz Rocha, José Carlos Vaz, and Diogo Cortiz. 2025. *Workshop Agentic Web: os impactos da IA na Web*. Relatório. Núcleo de Informação e Coordenação do Ponto BR (NIC.br), Centro de Estudos sobre Tecnologias Web – CEWEB.BR, São Paulo, Brazil. <https://acervo.ceweb.br/acervos/conteudo/d583800d-e2ec-47f6-bf09-fa6552eabc0a> ISBN 978-65-85417-99-0.
- [23] Chaojia Yu, Zihan Cheng, Hanwen Cui, Yishuo Gao, Zexu Luo, Yijin Wang, Hangbin Zheng, and Yong Zhao. 2025. A Survey on Agent Workflow – Status and Future. In *2025 8th International Conference on Artificial Intelligence and Big Data (ICAIBD)*. 770–781. doi:10.1109/ICAIBD64986.2025.11082076